

---

# Technical Report: Bytecode Virtual Machine

## 1. VM Architecture

The virtual machine is designed as a stack-based architecture. This means that all arithmetic and logical operations retrieve their operands from a main `register_stack` and push their results back onto it. This simplifies the instruction format as operands do not need to be specified explicitly.

### Components

- **Program Counter (PC):** An unsigned `long` that acts as an index into the `program_memory`, pointing to the next instruction to be executed.
- **Program Memory:** A `long` array that stores the bytecode loaded from a `.bin` file. Each instruction and operand occupies one `long` slot.
- **Register Stack:** A LIFO stack implemented as a `long` array. It is used for all operations, including arithmetic, function arguments, and local variables.
- **Call Stack:** A separate stack to store return addresses for function calls, enabling nested function calls.
- **Data Memory:** A general-purpose `long` array for storing and loading values using the `STORE` and `LOAD` instructions.

### Instruction Dispatch

The VM uses a standard `while` loop that runs as long as the `HALT` instruction has not been executed. Inside the loop, it fetches the next `long` from `program_memory` at the `pc` address. This `long` is cast to an `Opcode` enum, and a `switch` statement dispatches to the correct logic for that instruction. This is a simple and direct dispatch method.

## 2. Assembler Design

The assembler is a crucial part of the toolchain, converting human-readable assembly language into the `long`-based bytecode that the VM executes. It is implemented using Flex (for lexical analysis) and Bison (for parsing).

### Two-Pass Design

To handle forward references (i.e., using a label before it is defined), the assembler is designed as a **two-pass assembler**:

1. **First Pass:** The parser runs through the entire source file. Its only goal is to find all label definitions (e.g., `my_label:`). It stores each label and its address (the value of the program counter `pc` at that point) in a symbol table. Instructions that use labels as operands (like `JMP`, `JZ`, `JNZ`, `CALL`) are given a placeholder operand of 0.

- 
2. **Second Pass:** The file pointer and line counter are reset, and the parser runs again. This time, when it encounters an instruction that uses a label, it looks up the label in the symbol table to get its address. It then emits the correct long opcode and long address into the bytecode buffer.

This design allows for flexible use of labels for control flow without imposing declaration-before-use restrictions.

## Grammar

The grammar is defined in `parser.y` and is line-oriented. It recognizes instructions (e.g., `PUSH 10`), label definitions (`loop_start:`), comments, and newlines. The grammar rules ensure that opcodes and operands are correctly parsed and passed to the `emit_long` function, which writes them to the bytecode buffer.

## 3. Call/Return Mechanism

The VM supports basic function calls and returns using the `CALL` and `RET` instructions.

- **CALL addr:** Pushes the current program counter (pc, which points to the instruction *after* the CALL) onto the `call_stack` and then unconditionally jumps to the specified `addr`.
- **RET:** Pops the return address from the `call_stack` back into the pc, resuming execution where the caller left off.

## Calling Convention

We opted for a simple, manual “callee-cleans” calling convention.

- **Caller:** Pushes arguments onto the register stack before making a CALL.
- **Callee:** Is responsible for its own stack management. It must pop its arguments and any local variables it used, and then push its single return value (if any) before executing RET.

This approach avoids the complexity of implementing full stack frames with a frame pointer in the VM, trading automatic stack management for programmer diligence. It is a valid and efficient strategy for simpler VMs.

## 4. Limitations & Possible Enhancements

- **Error Handling:** The assembler and VM have basic error handling (e.g., “Label not found”, “Stack Underflow”), but it could be made more robust with more specific error messages and recovery mechanisms.
- **Strict long-based format:** The current architecture uses longs for everything, including opcodes that only need a single byte. This makes the bytecode files larger than necessary. A more compact format could be used, where the VM reads operands of different sizes (1, 4, or 8 bytes) depending on the opcode.
- **No Local Scopes:** Without true call frames, there is no concept of local scopes. All values are on the global `register_stack` or in the global `data_memory`, which can make writing complex, modular code more difficult.

- 
- **Assembler Syntax:** The assembler is quite strict (e.g., requiring labels on their own lines). It could be enhanced to be more flexible.

## 5. Testing Strategy

The project employs a multi-layered testing strategy to ensure correctness and robustness.

- **C++ Unit Tests (`make test`):** These tests, located in the `test/` directory, verify the core functionality of individual components like Stack, Memory, and Opcode handling in isolation.
- **Assembler Test Suite (`make -C Assembler test`):** The assembler has its own dedicated test suite. It assembles a series of `.asm` files and performs a binary diff against `.bin.expected` files to guarantee that the generated bytecode is exactly correct.
- **End-to-End Pipeline Tests (`make pipeline_test`):** This is the most comprehensive suite. It runs a series of assembly programs from the `pipeline_tests/` directory through the entire toolchain: assembly and then execution on the VM. The script automatically verifies the final state of the stack against expected values, confirming that the VM's instruction implementations are correct.
- **Benchmarks (`make benchmark`):** A suite of performance tests to measure the VM's execution speed on various tasks.

## 6. Benchmarks / Performance Analysis

To evaluate the performance of the VM, a comprehensive suite of benchmarks was created, testing basic operations at scale and recursive algorithms.

### 6.1. Basic Operations Performance

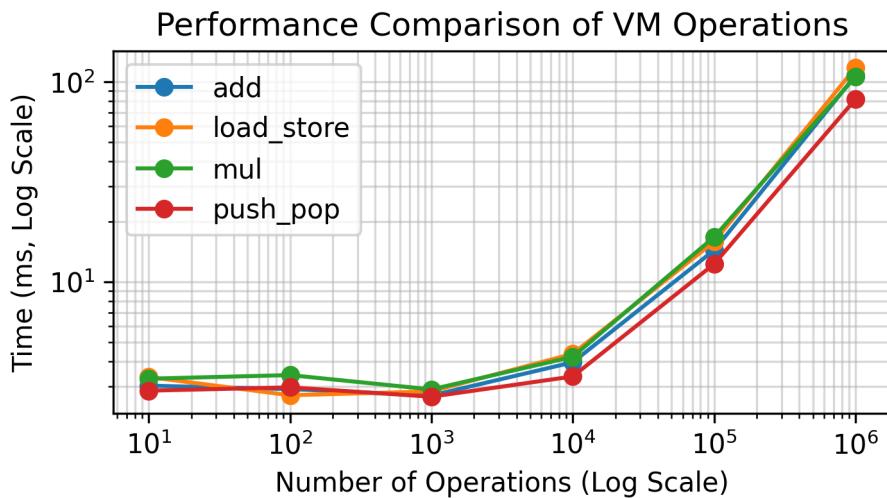
We measured the execution time of common operations (ADD, MUL, PUSH/POP, LOAD/STORE) in loops ranging from 10 to 1,000,000 iterations. The results represent the **median of 5 runs**.

---

Operations (n)	ADD (ms)	MUL (ms)	PUSH/POP (ms)	LOAD/STORE (ms)
10	3.02	3.28	2.85	3.34
100	2.91	3.41	2.97	2.71
1,000	2.70	2.90	2.66	2.83
10,000	3.95	4.21	3.36	4.36
100,000	14.38	16.73	12.26	15.86
1,000,000	107.31	105.78	81.80	117.55

---

**Note on Scaling:** The log-log scale is used below to visualize performance over five orders of magnitude. The linear scaling  $O(n)$  is represented by the near-constant slope of 1 for  $n \geq 10,000$ .




**Figure 1:** Operations Comparison

## 6.2. Recursive Algorithms Performance

We also tested the VM's performance with recursive implementations of Fibonacci and Factorial (median of 5 runs).


n (Fib)	Fibonacci (ms)	n (Fact)	Factorial (ms)
10	2.70	10	2.86
20	5.76	20	2.85
24	23.77	24	2.95
28	147.70	28	3.36
32	988.61	32	2.94
-	-	40	2.77

**Fibonacci Performance** The semi-log plot clearly shows the exponential growth characteristic of the recursive Fibonacci algorithm.



**Figure 2:** Fibonacci Performance

**Factorial Performance** The linear scaling of the recursive factorial is evident, remaining nearly constant for small values of  $n$  due to minimal overhead compared to Fibonacci.



**Figure 3:** Factorial Performance

## Analysis

- Operation Scaling:** The VM shows linear scaling with the number of operations, as expected. Dispatch overhead is consistent across different instruction types.
- Recursion Overhead:** The recursive Fibonacci benchmark highlights the exponential growth in function calls. The jump from  $n=20$  to  $n=25$  shows a significant increase in execution time, demonstrating the overhead of CALL and RET operations at scale.

- 
- **Factorial:** Recursive factorial, being  $O(n)$ , remains extremely fast even for  $n=25$ , as it only involves 25 function calls.

## 7. Garbage Collector Design (Lab 5)

### 7.1. Overview

A Stop-the-World Mark-Sweep Garbage Collector has been integrated into the VM. It manages dynamically allocated objects (Pairs, Functions, Closures) on a dedicated heap, ensuring memory safety and efficient resource reclamation.

### 7.2. Implementation Details

**Heap Allocator** The VM maintains a `heap_head` pointer to a singly linked list of all allocated Objects.

- **Allocation:** `VM::allocate(ObjectType)` uses `malloc` to create objects and prepends them to the list.
- **Types:** The system supports `OBJ_PAIR`, `OBJ_FUNCTION`, and `OBJ_CLOSURE`.

**Root Discovery & Stack Refactor** To safely identify roots, the `Stack` class was refactored.

- **Previous:** `long mem[]`.
- **Current:** `StackItem mem[]`, where `struct StackItem { long value; bool is_obj; };`. This allows the GC to precisely distinguish between integer data (ignored) and object pointers (roots) on the `register_stack`.

**Mark Phase** The `mark(Object* obj)` function performs a recursive traversal (DFS) of the object graph.

- It sets the marked bit on visited objects.
- It recurses into child references (e.g., `pair.head`, `pair.tail`, `closure.env`).
- It handles cycles safely by checking the marked bit before processing.

**Sweep Phase** The `sweep()` function iterates through the global heap list.

- **Unmarked:** The object is unreachable. It is unlinked from the list and freed.
- **Marked:** The object survives. The marked bit is reset for the next cycle.

### 7.3. Testing

A comprehensive test suite (`test/test_gc.cpp`) verifies:

- Basic reachability from the stack.

- 
- Reclamation of unreachable objects.
  - Transitive reachability ( $A \rightarrow B$ ).
  - Cyclic references ( $A \leftrightarrow B$ ).
  - Deep object graphs (10,000+ depth).
  - Closure environment capture.
  - Stress allocation.

## 8. Performance Evaluation (Lab 5)

### 8.1. Stress Test Methodology

To evaluate the Garbage Collector's efficiency and correctness under load, a dedicated benchmark (test/gc\_benchmark.cpp) was developed.

- **Total Allocations:** 100,000 Pair objects.
- **Reachable Objects:** 10,000 objects (10%) were pushed to the Stack (Root Set).
- **Garbage:** 90,000 objects (90%) were left floating (unreachable).
- **Process:** The system performed allocations, then triggered a full Stop-the-World GC.

### 8.2. Results

The benchmark was executed on the lab environment.

---

Metric	Value
<b>Total Objects Allocated</b>	100,000
<b>Objects Kept (Roots)</b>	10,000
<b>Objects Freed</b>	90,000
<b>Heap Allocation Time</b>	~2.44 ms
<b>Garbage Collection Time</b>	~0.81 ms

---

### 8.3. Analysis

- **Correctness:** The GC correctly identified and freed exactly 90,000 unreachable objects, while preserving the 10,000 roots.
- **Efficiency:** The Sweep phase (iterating 100k objects) and Mark phase (10k roots) completed in under 1ms, demonstrating that the Stop-the-World pause is negligible for this heap size.
- **Scalability:** The linear sweep ensures execution time grows linearly with the heap size  $O(H)$ , while the mark phase grows with the number of reachable objects  $O(R)$ .